

# A new NS3 Implementation of CCNx 1.0 Protocol

Marc Mosko\*, Ramesh Ayyagari, Priti Goel, Eric Holmberg, Mark Konezny

## Abstract

The `ccns3Sim` project is an open source implementation of the CCNx 1.0 protocols for the NS3 simulator. We describe the implementation and several important features including modularity and process delay simulation. The `ccns3Sim` implementation is a fresh NS3-specific implementation. Like NS3 itself, it uses C++98 standard, NS3 code style, NS3 smart pointers, NS3 xUnit, and integrates with the NS3 documentation and manual. A user or developer does not need to learn two systems. If one knows NS3, one should be able to get started with the CCNx code right away. A developer can easily use their own implementation of the layer 3 protocol, layer 4 protocol, forwarder, routing protocol, Pending Interest Table (PIT) or Forwarding Information Base (FIB) or Content Store (CS). A user may configure or specify a new implementation for any of these features at runtime in the simulation script. In this paper, we describe the software architecture and give examples of using the simulator. We evaluate the implementation with several example experiments on ICN caching.

## keywords

ICN, CCNx, NS3, Simulation

*Computing Science Laboratory, PARC*

\*corresponding author: marc.mosko@parc.com

## Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Node Architecture</b>	<b>2</b>
2.1	Modeling delays . . . . .	4
2.2	CCNxL3Protocol (Layer 3) . . . . .	5
2.3	CCNxStandardForwarder . . . . .	5
2.4	CCNxPortal (Layer 4) . . . . .	6
<b>3</b>	<b>Routing</b>	<b>6</b>
3.1	Installing a routing protocol . . . . .	7
3.2	Name Flooding Protocol Routing . . . . .	7
<b>4</b>	<b>Applications</b>	<b>8</b>
4.1	Consumer-Producer . . . . .	8
<b>5</b>	<b>Experiments</b>	<b>8</b>
<b>6</b>	<b>Conclusions</b>	<b>9</b>
	<b>References</b>	<b>9</b>

## 1. Introduction

The `ccns3Sim` [1] open source project is an NS3 [2] implementation of the CCNx 1.0 protocol specifications [8, 7]. It is released as a source code module for the NS3 simulator. It runs with an unmodified NS3 simulator, though one patch is required to add an Ethertype to the PPP implementation. `ccns3Sim` includes many examples and documentation.

The `ccns3Sim` implementation is a fresh NS3-specific implementation. Like NS3 itself, it uses C++98 standard, NS3 code style, NS3 smart pointers, NS3 xUnit, and integrates with the NS3 documentation and manual. A user or developer does not need to learn two systems. If one knows NS3, one should be able to get started with the CCNx code right away. A developer can easily use their own implementation of the layer 3 protocol, layer 4 protocol, forwarder, routing protocol, Pending Interest Table (PIT) or Forwarding Information Base (FIB) or Content Store (CS). A user may configure or specify a new implementation for any of these

features at runtime in the simulation script.

The general architectural model of `ccns3Sim` is to use an abstract base class as an *interface* and then provide an implementation class of that interface. For example, the `CCNxPit` abstract base class defines the API of the PIT table and the implementation class `CCNxStandardPit` realizes a PIT based on the current specifications. Because the abstract class inherits from `ns3::Object`, one can use run-time binding, making it easy for a user to substitute new implementations.

We introduce the Name Flooding Protocol (NFP), a simple distance-vector routing protocol for CCNx. It provides basic loop-free equal-cost multipath routing for CCNx names. Each *anchor* for a prefix (device advertising a prefix) uses its own sequence number for each advertisement, which are then ordered by distance. Each anchor is an independent successor graph for the prefix and will induce its own equal-cost multipath graph. If two or more anchors advertise the same prefix, NFP allows unequal cost multipath to each anchor.

The `ndnSim` project [6] is a customized NS3 distribution that uses the regular NDN open source code `ndn-cxx` (NDN C++ library with eXperimental eXtensions) and `nfd` (NDN Forwarding Daemon). This approach has a significant advantage in that it uses the same libraries as regular NDN applications and forwarders, so there is high code re-use and the simulator stays up-to-date with the non-simulation release. The disadvantage of this approach is that simulation must use the production code, it introduces two object models, two memory management models, and two coding style models. It may also make certain instrumentation that one would like to see in simulation difficult, because the `ndn-cxx` and `nfd` libraries were not designed with simulation in mind. For example, it can be difficult to simulate processing delays in data structures.

`ccnSim` [3] is a C++ simulator implemented in `Omnet++`. The authors claim it is a highly scalable simulator able to support content stores with up to  $10^6$  chunks and catalogs of up to  $10^8$  files. The current version, `ccnsim-v0.4`, may support much larger catalogs and operates up to 100x faster than the previous version. `ccnSim` supports several caching mechanisms (decision strategies), such as leave copy everywhere, leave copy down, and betweenness centrality. It also supports several cost and replacement functions to determine if a content is cached what what gets evicted. Our present work, `ccns3Sim`, is a more general purpose simulator for transport protocols, routing protocols, and other protocol features, it is not a highly optimized for content caching simulations. NS3's operational model of serializing and deserializing each packet at each hop means that content store copies are not memory references to a common repository. This makes the initial code release of `ccns3Sim` inefficient for large content caching experiments compared to `ccn-`

`Sim`.

The remainder of the paper is organized as follows. Section 2 describes `ccns3Sim`'s software architecture. It also describes the robust delay model and shows how to implement both input delays and processing delays. Section 2.2 describes the layer 3 components, and in particular the forwarder and associated tables. Section 2.4 describes the layer 4 components, which we call the `CCNxPortal` interface. Section 3 describes the routing protocol and the details of NFP. Section 4 describes the application model and the example consumer-producer app. Section 5 presents several experiments conducted with `ccns3Sim`. Section 6 concludes the paper and summarizes its contributions.

## 2. Node Architecture

`ccns3Sim` is designed such that a single `ns3::Node` may have zero or more CCNx applications running on it, plus a forwarder and routing protocol. An application running on a node uses one or more `CCNxPortal` interfaces to pass messages to the local forwarder. Each application instantiates its own Portals, similar to how an Internet application uses Sockets. A routing protocol is similar to an application in that it uses one or more Portals to communicate with peers, but it derives from `CCNxRoutingProtocol` instead of `CCNxApplication`. A routing protocol also uses the `CCNxForwarder` interface on a node to manipulate the FIB inside the forwarder. `ccns3Sim` can operate with any NS3 `NetDevice`. We have tested with point-to-point, CSMA, WiFi, and virtual devices using UDP tunnels.

The CCNx NS3 architecture is shown in Fig. 1. Every major component is described by an interface (abstract base class) and inherits from the `ns3::Object`. This allows run-time late binding of each piece to a user-specified implementation in the simulation script. The core of the architecture is the `CCNxL3Protocol`, which is the switching fabric for the layer 3 protocol. It is connected to network devices below and layer 4 protocols above. There is also a special interface `CCNxForwarder` to the forwarding engine, which is responsible for selecting the egress connection of each input packet. We describe each of these components in more detail below.

The CCNx node architecture differs from the NS3 Internet model in that we have separated the routing protocol from the forwarding process. In the NS3 model, the routing protocol, in addition to exchanging routing messages with peers and executing a routing process, also handles calls to `RouteInput()` and `RouteOutput()`. In the CCNx architecture, the routing protocol only exchanges messages with peers and executes a routing process. It maintains a Routing Information Base (RIB). Using the `CCNxForwarder` interface, the routing protocol calls `AddRoute()` and `RemoveRoute()` to manage the Forwarding Information Base (FIB) in-



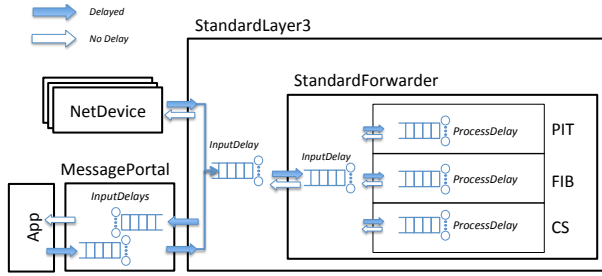


Figure 3. Layer delays

for example, taken PCAP traces from NS3 and replayed them on real networks to our software forwarders. The current NS3 implementation is only partially configurable for modular codecs, but we expect to have fully configurable codecs in the near future. Because CCNx 1.0 uses the hierarchical nesting of TLV types to determine containment, we will use a notation like SNMP MIB OIDs to specify codecs. For example, TLV type 1 is Interest message and TLV type 0 is its CCNxName, so “.1.0” would specify the codec for the name. TLV type 1 is the Payload of the Interest, so “.1.1” is used to lookup the payload’s codec, and so on. This model allows a user to either replace a specific codec or introduce new TLV types and new codecs at run time. A codec may be used by multiple OIDs. For example “.1.0” and “.2.0” are both CCNxName codecs. The first is for an Interest message and the second is for a Content Object message.

## 2.1 Modeling delays

The `CCNxDelayQueue` is a general queuing class to delay events. Several of the major CCNx forwarding components have delay queues to model computation time. We use the terminology *input delay* to mean a delay that only depends on the input to the delay queue, not the processing done. The term *processing delay* means a delay that depends on the processing, such as the number of tables consulted or the number of name components searched. Both use the same data structure, it is only a difference of where the processing is done.

We provide examples of using both input delay and processing delay models. Input delay models are useful for layers where the delay only depends on the item being delayed. One example is the delay in `CCNxStandardL3`, which is only a fixed delay. Another example is `CCNxStandardPit`, which uses a linear function of the CCNxName byte length. Processing delay models are useful for cases where the delay depends on the result of the processing. `CCNxStandardFib`, for example, uses a linear function of the number of FIB lookups done to find the longest matching prefix. `CCNxStandardContentStore` uses a constant delay if no match is found and a linear delay in terms of the size of the Content Object matched otherwise.

A `CCNxDelayQueue<T>` is constructed with the




Figure 4. Delay Queue

number of parallel servers of the queue (which enforces the delay) and two functions. The first function `GetServiceTime(T item)` is called whenever a queue item is moved to the head-of-line and there is a free server for it. The second function `DequeueCallback(T item)` is called after the service time is over. It dequeues the item back to the user’s code.

```

1 Time GetServiceTime(T item) {
2     return MicroSeconds(2);
3 }
4 void DequeueCallback(T item) {
5     MyResultClass result = lookup(item);
6     callNextStage(item, result);
7 }

```

Listing 2. Modeling input delay

When used to model input delay, the `GetServiceTime(T item)` function only depends on the input to the function. It could be a constant, or some function of the queue item. The delay, for example, could be linear in the number of bytes of the packet or some function of the number of name components. An example of a constant delay using this model is in Listing 2. Once the service time is over, the `DequeueCallback(T item)` function carries on processing the queue item.

```

1 Time GetServiceTime(T item) {
2     MyResultClass result = lookup(item)
3     item->SetResult(result);
4     return result->GetServiceTime();
5 }
6
7 void DequeueCallback(T item) {
8     callNextStage(item);
9 }

```

Listing 3. Modeling processing delay

When used to model processing delay, the delay calculation depends on the processing function. A FIB lookup delay, for example, might depend on the number of name components and the length of each component. The typical usage pattern is for `GetServiceTime(T item)` to perform the processing on `item` and put the result back in `item` (e.g. `item->SetResult(...)`) and then return the amount of simulation delay to wait.



Once the delay is over and the item is passed to `DequeueCallback(T item)`, no further calculation is done and `item` is passed to the next step in the service pipeline. Listing 3 shows an example of this usage.

## 2.2 CCNxL3Protocol (Layer 3)

As shown in Figure 1, the `CCNxL3Protocol` interface (abstract base class), implemented by `CCNxStandardLayer3` class, is the central glue between the network and the node. Besides the necessary APIs to send and receive packets, it also performs these functions:

- *Layer 3 Interfaces*: It maintains a layer 3 abstraction of network devices via the `CCNxL3Interface` class. This maintains layer 3 data about an interface, such as if CCNx is forwarding on it or if the interface is administratively up or down. It also tracks how to broadcast on a network device.
- *Connection Table*: The connection table stores information about each adjacency using the `CCNxConnection` abstract base class. It is implemented by `CCNxConnectionDevice` for layer 2 devices and `CCNxConnectionL4` for layer 4 transport connections. Each individual peer on a network interface has its own connection. Each layer 2 interface that supports broadcast has a special connection in the connection table for its broadcast address.
- *Prefixes*: Registering a prefix from Layer 4 means the FIB will be updated to add to a specific `CCNxConnectionL4` for the prefix. Unregistering a prefix removes that connection from the FIB.
- *Anchors*: Anchors are like registering and unregistering a prefix, but in addition the routing protocol will be notified that the local node is an anchor. The routing protocol will then begin advertising the prefix as locally reachable.

```
1 AcmeLayer3Helper layer3;
2
3 CCNxStackHelper ccnx;
4 ccnx.SetLayer3Helper(layer3);
5 ccnx.Install (nodes);
```

**Listing 4.** Using a custom Forwarder

The `CCNxL3Protocol` is replaceable via the simulation script. A user could create, for example, the `AcmeLayer3` that implements the `CCNxL3Protocol` interface and the `AcmeLayer3Helper` that implements the `CCNxLayer3Helper` abstract base class. A user may then instantiate the helper in the simulation script and pass it to `CCNxStackHelper::SetLayer3Helper()`, as shown in Listing 4.

## 2.3 CCNxStandardForwarder

The `CCNxForwarder` abstract base class represents the interface to a forwarder. A forwarder manages all the state necessary to find the next hop for packet. In a standard implementation, this means it has a FIB to forward Interest and a PIT to forward Content Objects. The class `CCNxStandardForwarder` implements the standards-based forwarder. The principle function of the forwarder is to service calls to `RouteInput()` and `RouteOutput()`. These functions, whose names are taken from the NS3 Internet module, are the interface to forward a packet. `CCNxL3Protocol` calls the first when it receives a packet from a network device and calls the second when it receives a packet from a Layer 4 protocol. The result of those functions is a vector of `CCNxConnections` on which forward the packet.

A user may configure the forwarder via the `CCNxStandardForwarderHelper`, which implements the `CCNxForwarderHelper` abstract base class. A user may entirely replace the forwarder by implementing their own forwarder and forwarder helper. For example, as shown in Listing 5, the class `acme::AcmeForwarder` implements the `CCNxForwarder` interface and is configured via `acme::AcmeForwarderHelper`.

```
1 AcmeForwarderHelper forwarder;
2 forwarder.UseTwoCopy(true);
3
4 CCNxStackHelper ccnx;
5 ccnx.SetForwardingHelper(forwarder);
6 ccnx.Install (nodes);
```

**Listing 5.** Using a custom forwarder

The `CCNxStandardForwarderHelper` exposes three factories to allow a user to customize or replace the three main functional components: PIT, FIB, and ContentStore. This abstraction is specific to the `CCNxStandardForwarderHelper`, not the `CCNxForwardingHelper` because those tables may not exist in other forwarders, such as a label-swapping forwarder.

The tables are represented by an `ns3::ObjectFactory`, not a helper. This is because, in ns3 usage, a helper not only creates an `ns3::Object`, it also aggregates it to a node. The factories for the forwarder tables do not perform an aggregation step; they are a pure factory, so we call them a factory not a helper. Suppose you have `AcmePit` that implements the `CCNxPit` interface. One may use code similar to Listing 6 to replace the PIT used by the standard forwarder.

Similar to the PIT, `CCNxStandardForwarderHelper` takes an `ObjectFactory` for the FIB and Content Store. A user may provide their own implementations by creating their own factory and passing it to the `CCNxStandardForwarderHelper`.

```

1 AcmePitFactory acmePitFactory;
2 acmePitFactory.SetMaxEntries(1000);
3
4 CCNxStandardForwarderHelper forwarder;
5 forwarder.SetPitFactory(acmePitFactory);
6
7 CCNxStackHelper ccnx;
8 ccnx.SetForwardingHelper(forwarder);
9 ccnx.Install (nodes);

```

Listing 6. Using a custom PIT

## 2.4 CCNxPortal (Layer 4)

The `ccns3Sim` layer 4 is the `CCNxPortal`. It is modeled after the `ns3::Socket`, which has specializations for UDP and TCP. Likewise, Portal has a specialization `CCNxMessagePortal`, which is a simple Interest and Content Object passing portal that does no processing on the messages. Future specializations are a Manifest portal and a Chunked portal that provides reliable, in-order congestion control.

```

1 ObjectFactory f("ns3::ccnx::↵
    CCNxMessagePortalFactory");
2 Ptr<Object> p = f.Create<Object> ();
3 node->AggregateObject (p);

```

Listing 7. Binding a Portal Factory

```

1 class Tempor
2 {
3 private:
4     Ptr<CCNxPortal> m_portal;
5
6     void ReceiveNotify(Ptr<CCNxPortal> p) {
7         Ptr<CCNxPacket> packet;
8         while ((packet = p->Recv ())) {
9             // use packet
10        }
11    }
12 public:
13     void Labore(Ptr<Node> node) {
14         m_portal = CCNxPortal::CreatePortal (node, ↵
            TypeId::LookupByName ("ns3::ccnx::↵
                CCNxMessagePortalFactory"));
15
16         m_portal->SetRecvCallback (MakeCallback (&↵
            Tempor::ReceiveNotify, this));
17    }
18 };

```

Listing 8. Using a Portal

Similar to `ns3::Socket`, when `CCNx` is instantiated on a node via `CCNxStackHelper`, it creates a `CCNxPortalFactory` for each portal type (in this case just a message portal) and aggregates it to the node. This allows applications to ask for the portal factory by name and create an instance of `CCNxPortal` with the desired implementation. This method also allows the user to aggregate their own factories from the simulation script and use those portals in the same way.

```

1 // Local forwarder only
2 virtual bool RegisterPrefix (Ptr<const CCNxName>↵
    prefix) = 0;
3 virtual void UnregisterPrefix (Ptr<const ↵
    CCNxName> prefix) = 0;
4
5 // Local forwarder and routing protocol
6 virtual bool RegisterAnchor (Ptr<const CCNxName>↵
    prefix) = 0;
7 virtual void UnregisterAnchor (Ptr<const ↵
    CCNxName> prefix) = 0;
8
9 // SendTo() should change to take CCNxConnection
10 virtual bool Send (Ptr<CCNxPacket> packet) = 0;
11 virtual bool SendTo (Ptr<CCNxPacket> packet, ↵
    uint32_t connid) = 0;
12
13 virtual Ptr<CCNxPacket> Recv (void) = 0;
14 virtual Ptr<CCNxPacket> RecvFrom (Ptr<↵
    CCNxConnection> & incoming) = 0;

```

Listing 9. CCNxPortal data API

Listing 8 illustrates a common way to use a Portal inside a class. The class defines a private variable `m_portal` that is instantiated in the method `Labore`. After the portal is created from its portal factory, the method must also set the data receive callback, which in this case is `Tempor::ReceiveNotify()`. This design is similar to how the NS3 `Socket` class works. When a packet arrives in a portal, it will call the receive notify method to inform the owner of the portal that data is ready. The owner of the portal may begin reading immediately (in the callback) or could defer reading to a later time.

## 3. Routing

`ccns3Sim` provides two ways to do routing. The first is to use the `CCNxStaticRoutingHelper` to create static routes between two nodes. The second is to use the NFP routing protocol (Section 3.2), which is dynamic routing protocol. A user may also provide their own routing protocol by implementing the `CCNxRoutingProtocol` interface.

A routing protocol uses one or more `CCNxPortals` to send and receive packets from peers. It may also use `CCNxPortal::SendTo()` to send a packet to a specific connection or rely on the FIB. A routing protocol uses the `CCNxForwarder` interface to call `AddRoute()` and `RemoveRoute()` on the forwarder, which allows a very loose coupling between the routing protocol and the forwarder.

A user specifies which routing protocol to use as shown in Listing 10. As with other components, the user instantiates a routing helper, such as `NfpRoutingHelper`, sets any desired parameters on it, then passes the helper to the `CCNxStackHelper`.

### 3.1 Installing a routing protocol

Instantiating a routing protocol uses the `NfpRoutingHelper`, as shown in Listing 10. Once the script sets optional parameters, such as the `HelloInterval`, it then adds it to the `CCNxStackHelper` before installing the stack helper on nodes.

```
1 NfpRoutingHelper nfp;
2 nfp.Set("HelloInterval", TimeValue(Seconds (1)))←
3 ;
4 CCNxStackHelper ccnx;
5 ccnx.SetRoutingHelper(nfp);
6 ccnx.Install (nodes);
```

**Listing 10.** Installing a routing protocol

### 3.2 Name Flooding Protocol Routing

The Name Flooding Protocol (NFP) is a distance vector routing protocol for CCNx names. It creates independent directed acyclic graphs (DAGs) for each pair (*anchorName*, *prefix*). An anchor name is a unique identifier for a node that advertises a prefix. We use a CCNxName for both the *anchorName* and *prefix*. The *prefix* is a CCNx name *prefix* for matching against Interest messages to determine where to send them. This construction allows each node to keep the maximum information about *prefix* reachability at the expense of extra signaling and storage costs. Node should still forward only towards the least cost anchor, as there is no guarantee that two anchor paths will not cause a loop for the same *prefix*. NFP is intended as a demonstration of a CCNx routing protocol, and exhibits worse performance than other protocols, such as in [5].

NFP advertises the tuple (*prefix*, *anchorName*, *anchorPrefixSeqnum*, *distance*). Each (*prefix*, *anchorName*) pair is considered independently from other advertisements. This means that NFP maintains routes to all anchors that advertise a specific name *prefix*. Each anchor is responsible for incrementing its own *anchorPrefixSeqnum* with each new advertisement. As the advertisement travels over the network, its distance is always increasing with each hop. We do not require min-hop link costs, it could be any positive distance per hop.

Each anchor has a unique name, for example a cryptographic hash of a public key. In our experiments, we use a name that is a 32-byte value, similar to what a SHA-256 hash of a public key would be.

An advertisement for (*prefix*, *anchorName*) is considered feasible if the *anchorPrefixSeqnum* is greater than any previously seen sequence number, or *anchorPrefixSeqnum* being equal to a previous advertisement and the distance is no greater than any previous advertisement seen at the node. An advertisement is *strictly better* if the sequence number is larger or begin equal the distance is strictly less. A withdraw message is feasible if

its sequence number is larger than the currently known sequence number and it came from a current predecessor for the (*prefix*, *anchorName*) pair. A withdraw has no distance associated with it. This feasibility criteria allows equal-cost multipath to each anchor. Periodically (e.g. every 10 seconds) the anchor advertises its *prefixes*. Each link (neighbor adjacency) has a positive cost. We use “1” for all link costs.

Each routing node will accept only feasible advertisements. When a node receives a strictly better advertisement, it erases all prior predecessors for the (*prefix*, *anchorName*) pair and uses only the strictly better advertisement. It immediately (with jitter) forwards the advertisement with increased cost to its peers. When a node receives an equal advertisement, it adds the predecessor to the list of equal cost multi-paths for the (*prefix*, *anchorName*) pair and does not forward the advertisement. If a withdraw message changes the state of a (*prefix*, *anchorName*) pair to unreachable, it is forwarded immediately (with jitter). If it does not change the state, its only effect is to remove a specific predecessor from the set of multi-path predecessors.

NFP uses two routing messages: an advertisement and a withdraw. If a node loses reachability for a route, it may immediately send a withdraw message to its neighbors. Routes use soft state and will timeout after 3x the advertisement period. Advertisements and withdraws are sent in a routing message, which is in turn carried in the payload of a Interest message. Each routing message has a 1-hop *MessageSeqnum* specific to that 1-hop message. A message is only accepted if (*RouterName*, *MessageSeqnum*) is greater than the previously heard (*RouterName*, *MessageSeqnum*). A routing message is also an implicit Hello from that neighbor. It may be sent with no advertisements or withdraws to function simply as a Hello.

The neighbor table tracks Hello status for each neighbor. It has a callback to `NeighborStateChanged()` whenever the state of a neighbor changes (UP, DOWN, DEAD). We use a 3-state model for neighbors so we do not erase their *messageSeqnum* right away when they go DOWN. They have to timeout a second time to DEAD state before we erase the record. If a neighbor goes to DOWN state, we erase all next hops that use that neighbor. That may cause `PrefixStateChanged()` to be called if a (*prefix*, *anchorName*) becomes unreachable. On a DEAD change, we simply erase the record.

The *prefix* table stores the best advertisements we have seen for each (*prefix*, *anchorName*) pair. It has a callback to `PrefixStateChanged()` when a (*prefix*, *anchorName*) becomes reachable or unreachable. There currently is not a method to prune DEAD records like there is for neighbors.

## 4. Applications

CCNx applications inherit from the `CCNxApplication` base class, which in turn inherits from `ns3::Application`. This means that CCNx applications can use the `ns3::ApplicationContainer`, which provides easy mechanisms to start and stop applications on a node. We provide two example applications: a consumer and a producer.

### 4.1 Consumer-Producer

The consumer-producer example applications use a `CCNxRepository`. The repository is a name prefix plus a fixed number of content objects all of a fixed size. One repository may be shared between all producers (replicas), so there is no duplication of state. Each `CCNxProducer` will register as an anchor for the repository prefix (and thus advertise it via a routing protocol) and serve the repository's content via a `CCNxMessagePortal`. Each `CCNxConsumer` will request names from the repository with a fixed request rate. For each request, it uses `CCNxRepository::GetRandomName()` to pick a name from the repository. Currently, the repository only supports a uniform name distribution. A node may have zero or more producers on it.

```
1 Ptr<const CCNxName> prefix = Create<CCNxName> (<
  ("ccnx:/name=prefix");
2 uint32_t bytes = 124;
3 uint32_t count = 1000;
4 Ptr<CCNxContentRepository> repo = Create<
  CCNxContentRepository> (prefix, bytes, count);
```


**Listing 11.** Creating a repository

Listing 11 shows how to create a content repository with a given prefix. In this example, each content object, which will have its own individual name under the name `ccnx:/name=prefix`, will be 124 bytes of payload, and there will be 1,000 of them. The `Ptr<>` reference to the repo is then shared between producer nodes and consumer nodes (see the next listing).

Listing 12 shows how to create a set of consumer nodes for a repository `repo`. Because the `CCNxApplication` class follows the standard NS3 application model, we can use the `ApplicationContainer` to schedule starting and stopping an application. The same process applies to the `CCNxProducer`.

```
1 CCNxConsumerHelper consumer (repo);
2 consumer.SetAttribute ("RequestInterval", <
  TimeValue (Milliseconds (5)));
3 ApplicationContainer consumerApps = consumer.<
  Install (consumerNodes);
4 consumerApps.Start (Seconds (2));
5 consumerApps.Stop (Seconds (12));
```

**Listing 12.** Configuring a consumer



**Figure 5.** Computational Cost after Link Failure

## 5. Experiments

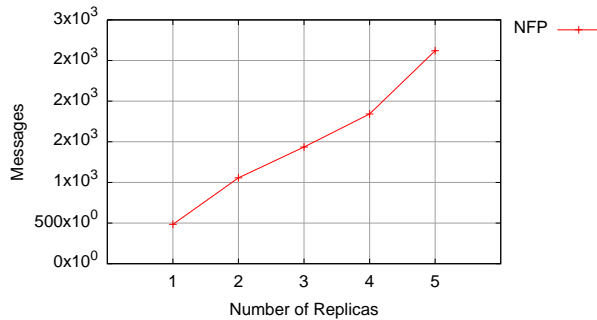
This section presents results of some experiments with the NFP routing protocol (see Section 3.2). The purpose is to illustrate using the simulator on a non-trivial topology with real routing traffic. These experiments are not a thorough evaluation of the NFP routing protocol, though they do show some of its characteristics. As mentioned in the prior section on NFP, because it treats each pair (*anchorName*, *prefix*) as an independent routing process, it will scale worse than LSCR [5] or DCR [4], which only advertise the best anchor for a prefix.

We used a similar methodology as in [5], which presents the routing protocol LSCR. The simulations are done on the 154 node AT&T topology using the simulator's NFP protocol. Because NFP is a distance vector protocol, the comparison is not exactly apples to apples. Instead of measuring link state advertisements (LSAs), we measure advertisement and withdraw routing update messages. We measure computational complexity the same: it is each call for an event (e.g. a timer or table operation) and each iteration through a loop control structure. We have not implemented the NLSR or LSCR protocols in our simulator, so we cannot directly compare results. The methodology in [5] uses 30 random nodes as anchors and randomly distributed 210 name prefixes among them. We repeat each experiment 20 times.

Figure 5 shows the computation cost associated with a single link failure. After an initial period (30 seconds), a single link is selected and failed by introducing a 100% packet loss. After 3 seconds, the Hello protocol times out and the peers on the link remove each other. Because the topology has a few nodes of high degree and many nodes of low degree, this usually results in one or a small number of nodes being disconnected. The computational cost is almost linear in the number of replicas because each (*anchorName*, *prefix*) pair is propagated, so when a leaf node has its link cut, it will lose a number of routes proportional to the number of replicas.

Figure 6 shows the number of routing update messages triggered by a link failure. As with the computa-





**Figure 6.** Routing messages after Link Failure

tional cost, the number of *withdraw* messages will scale linearly with the number of replicas.

In terms of processing efficiency, *ccns3Sim* executes these 50 second virtual time examples in under 40.1 seconds real time in a simulation with over 18.8M computational complexity and 221,880 packets. That comes out to an average execution time of 2.1 usec per routing protocol event on a MacBook Pro laptop with a 2.6 GHz core i7 CPU. The simulations use under 285 MB of RAM (just under 1.8 MB per node). Most of the memory is going to the content repositories.

## 6. Conclusions

The *ccns3Sim* module for NS3 implements the CCNx 1.0 protocols. It is a native NS3 implementation in C++98 that uses the NS3 coding style, memory and object management functions, and application style. Using modular NS3 helpers, a user can substitute or configure all the major CCNx components: the Layer 3 implementation, the forwarder, the routing protocol, the PIT, FIB, and Content Store with in the forwarder, and the layer 4 Portal. Each of these pieces includes a layer delay model to allow simulation of data structure and processing delays. We also provide an example routing protocol, Name Flooding Protocol (NFP) to illustrate how a dynamic routing protocol operates in *ccns3Sim*.

There is still more work to be done. Future work will expand the cryptographic side of the implementation to allow modeling the use of public key cryptography. The packet encoding currently does not support virtual payloads, which are a common simulation tool to reduce processing time. *ccns3Sim* is also missing Interest NACKs, organizational unit TLVs, and nameless objects. We also plan on incorporating the CCNx model within NS3's LTE model to simulate its use in mobile networks.

In terms of protocol implementation, future work will include several important features. The consumer-producer applications will support Zipf request distributions and the content store will support different placement (where to cache) and replacement strategies. There are also many optimizations to make content caching

simulations more efficient. We will also implement several different routing protocols and new transport protocols, including a chunk based reliable transport and a manifest based reliable transport.

## References

- [1] *ccns3Sim* open source project. <https://github.com/PARC/ccns3Sim/wiki>. Accessed: 2016-05-05.
- [2] NS3 simulator. <http://nsnam.org>. Accessed: 2016-05-05.
- [3] R. Chiocchetti, D. Rossi, and G. Rossini. *ccnsim*: An highly scalable ccn simulator. In *2013 IEEE International Conference on Communications (ICC)*, pages 2309–2314, June 2013.
- [4] J. Garcia-Luna-Aceves. Name-based content routing in information centric networks using distance information. In *Proceedings of the 1st International Conference on Information-centric Networking, ICN '14*, pages 7–16, New York, NY, USA, 2014. ACM.
- [5] E. Hemmati and J. Garcia-Luna-Aceves. A new approach to name-based link-state routing for information-centric networks. In *Proceedings of the 2nd International Conference on Information-Centric Networking*, pages 29–38. ACM, 2015.
- [6] S. Mastorakis, A. Afanasyev, I. Moiseenko, and L. Zhang. *ndnsim 2.0: A new version of the ndn simulator for ns-3*. Technical report, Technical Report NDN-0028, NDN, 2015.
- [7] M. Mosko, I. Solis, and C. Wood. Ccnx messages in tlv format. Internet-Draft draft-irtf-icnrg-ccnxmessages-02, IETF Secretariat, April 2016. <http://www.ietf.org/internet-drafts/draft-irtf-icnrg-ccnxmessages-02.txt>.
- [8] M. Mosko, I. Solis, and C. Wood. Ccnx semantics. Internet-Draft draft-irtf-icnrg-ccnxsemantics-02, IETF Secretariat, April 2016. <http://www.ietf.org/internet-drafts/draft-irtf-icnrg-ccnxsemantics-02.txt>.